

ALMA MATER STUDIORUM – UNIVERSITÀ DI BOLOGNA  
CAMPUS DI CESENA

---

Scuola di Ingegneria e Architettura  
Laurea Magistrale in Ingegneria e Scienze Informatiche

## USED CARS ANALYSIS

*Elaborato in*  
BIG DATA

*Presentata da*  
GABRIELE GUERRINI

---

Anno Accademico 2019 – 2020



# Table of contents

<b>1</b>	<b>Dataset and query</b>	<b>1</b>
1.1	Dataset . . . . .	1
1.2	Query . . . . .	1
1.2.1	Description . . . . .	1
1.2.2	Workflow . . . . .	1
<b>2</b>	<b>MapReduce</b>	<b>3</b>
2.1	Jobs . . . . .	3
2.1.1	Job 1: Preprocessing . . . . .	3
2.1.2	Job 2a: Opi . . . . .	4
2.1.3	Job 2b: Region . . . . .	5
2.1.4	Job 3: Join . . . . .	6
2.2	Performance evaluation . . . . .	7
<b>3</b>	<b>Spark</b>	<b>9</b>
3.1	Job description . . . . .	9
3.2	Performance evaluation . . . . .	14
<b>4</b>	<b>Conclusions</b>	<b>17</b>



# Chapter 1

## Dataset and query

### 1.1 Dataset

The dataset can be downloaded at:

<https://www.kaggle.com/austinreese/craigslist-carstrucks-data>.

It contains data about more than 400k used cars for sale in US. Each car is described by several info such as price, odometer, manufacturer, year, sale region and so on.

### 1.2 Query

#### 1.2.1 Description

The query to be executed is the following one:

“For each region, it must be found the OPI of the most widespread brand in such region, considering cars that use gas fuel only.

OPI is an acronym for Odometer-Price Index and represents the average ratio odometer/price. It is calculated upon all cars of a given brand in the country, regardless of other car features (fuel type, number of cylinders...).”

#### 1.2.2 Workflow

Figure 1.1 shows the adopted workflow<sup>1</sup>.

---

<sup>1</sup>It is not the optimal one since the same query could be executed in less than four job. The idea is to structure the workflow this way so that few different algorithms (e.g. filtering, projection, summarization, join) and optimizations (e.g. caching) can be applied.

The pipeline is made up of four different jobs, that is:

- **Job 1: Preprocessing:** The job executes all preprocessing operation needed to correctly use the dataset in next jobs. Firstly, it cleans the raw dataset by dropping records that contain missing values and by eliminating useless columns. Then, it executes the dataset growth operation, namely, each remaining record is copied by a certain replication factor so that the output dataset contains enough amount of data. The default replication factor is 100 (but it can be customized).
- **Job 2a: Opi:** It calculates the opi for each brand.
- **Job 2b: Region:** It calculates the most widespread brand for each region.
- **Job 3: Join:** It merges partial results of jobs 2a and 2b.

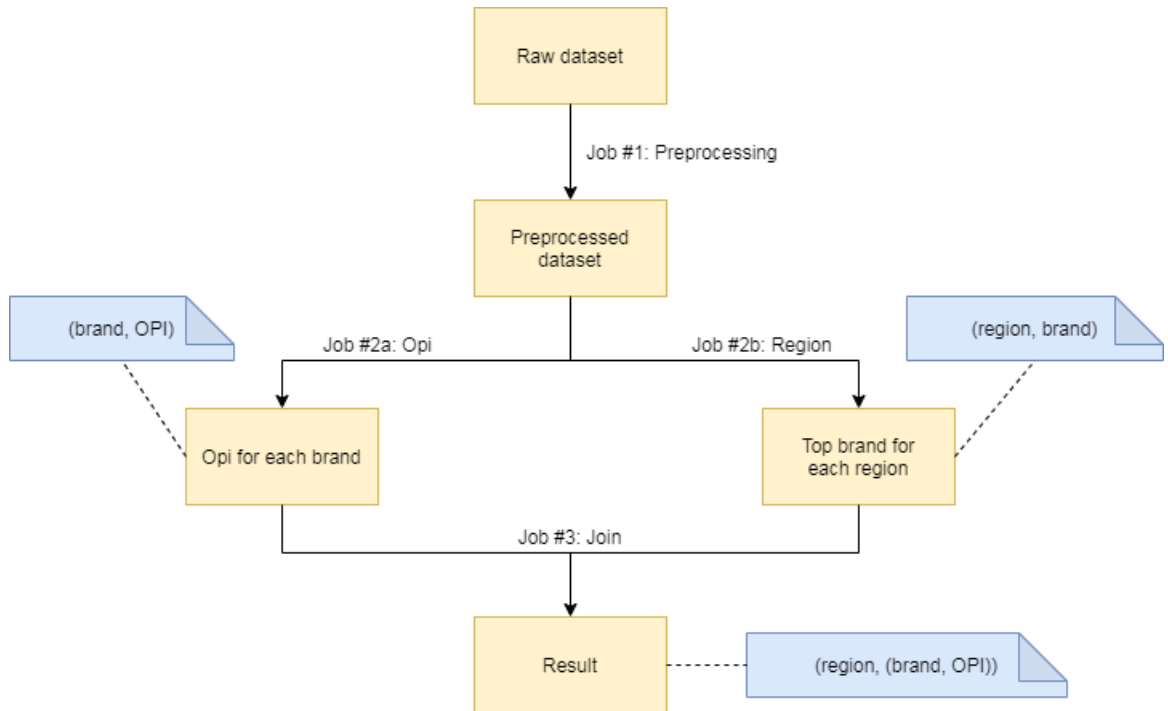


Figure 1.1: Adopted workflow for MR jobs.

# Chapter 2

## MapReduce

### 2.1 Jobs

#### 2.1.1 Job 1: Preprocessing

Job execution and interfaces are described in figure 2.1.

Raw records are read and parsed into “Car”, i.e. custom “Writable” objects, during map stage. Each car stores data about:

- Region
- Price
- Brand
- Fuel
- Odometer

The map output is a pair where the key is the default key used by Hadoop when reading text files and the value is the car itself.

The reduce stage replaces the default key with a “NullableWritable” so that the whole job output is a set of records yet and then publishes each record several times as specified by the replication factor.

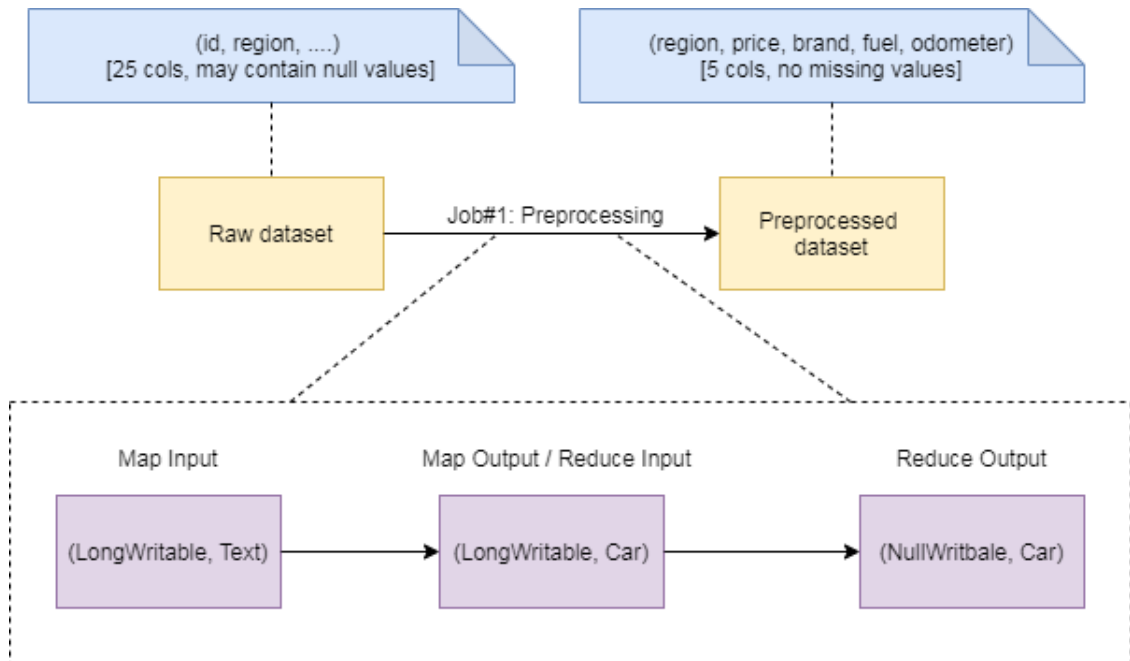


Figure 2.1: Description of “job 1: preprocessing” workflow and interfaces.

### 2.1.2 Job 2a: Opi

Job execution and interfaces are described in figure 2.2.

The map stage input value is a car encoded as text. The key is the Hadoop default one for text files and its irrelevant. The mapper calculates the OPI for such car using price and odometer fields, and it gives as output a pair consisting of (brand, OPI).

The reduce stage just calculates the average value using single OPI values and reports it as output.



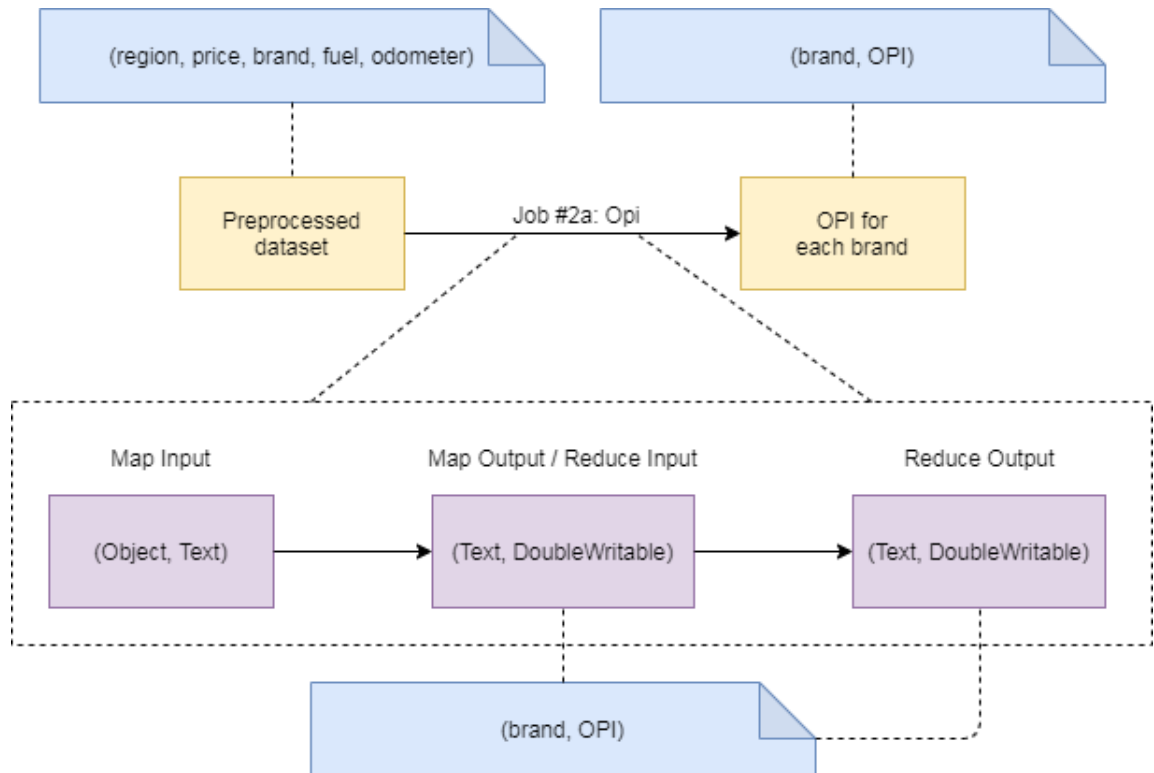


Figure 2.2: Description of "job 2a: opi" workflow and interfaces.

### 2.1.3 Job 2b: Region

Job execution and interfaces are described in figure 2.3.

The map stage input value is a car encoded as text. The key is the Hadoop default one for text files and its irrelevant. If the fuel type is correct, the mapper generates as output a pair (region, brand) so that reduce stage can operate per region basis.

The reducer, given a region, firstly calculates the cardinality for each brand by collecting the input pairs. Then, it calculates the maximum and gives it as output in the form of (region, brand).

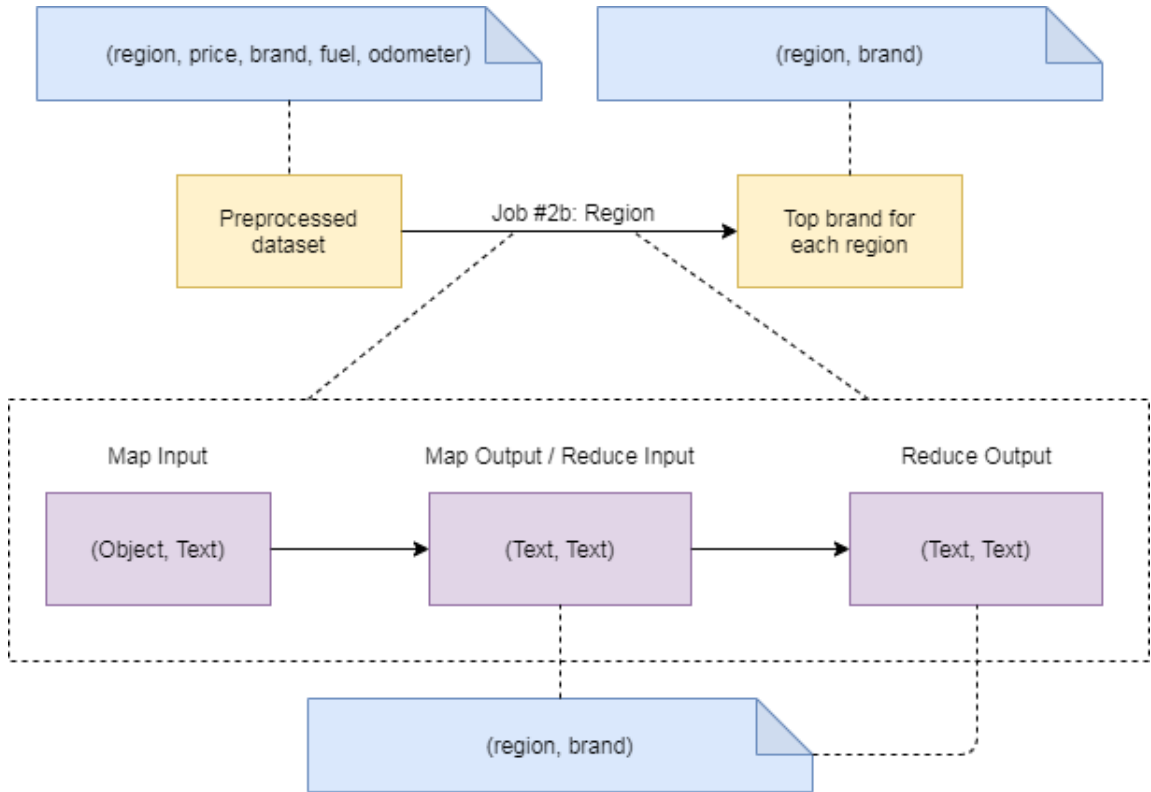


Figure 2.3: Description of “job 2b: region” workflow and interfaces.

### 2.1.4 Job 3: Join

Job execution and interfaces are described in figure 2.4.

The map stage takes as input pairs structured as  $(\text{Text}, \text{Text})$ . The source of a given pair can be from job 2a output or from job 2b output. By examining it, it is discovered the source and the pair is marked with a flag. Since the join key will be the brand, the output of map stage is a pair  $(\text{brand}, (\text{flag}, \text{region}))$ , where flag and region are encapsulated into a “JoinPair” (a custom “Writable”).

The reducer executes the standard join algorithm on input items and creates output pairs in the form  $(\text{region}, (\text{brand}, \text{OPI}))$ .

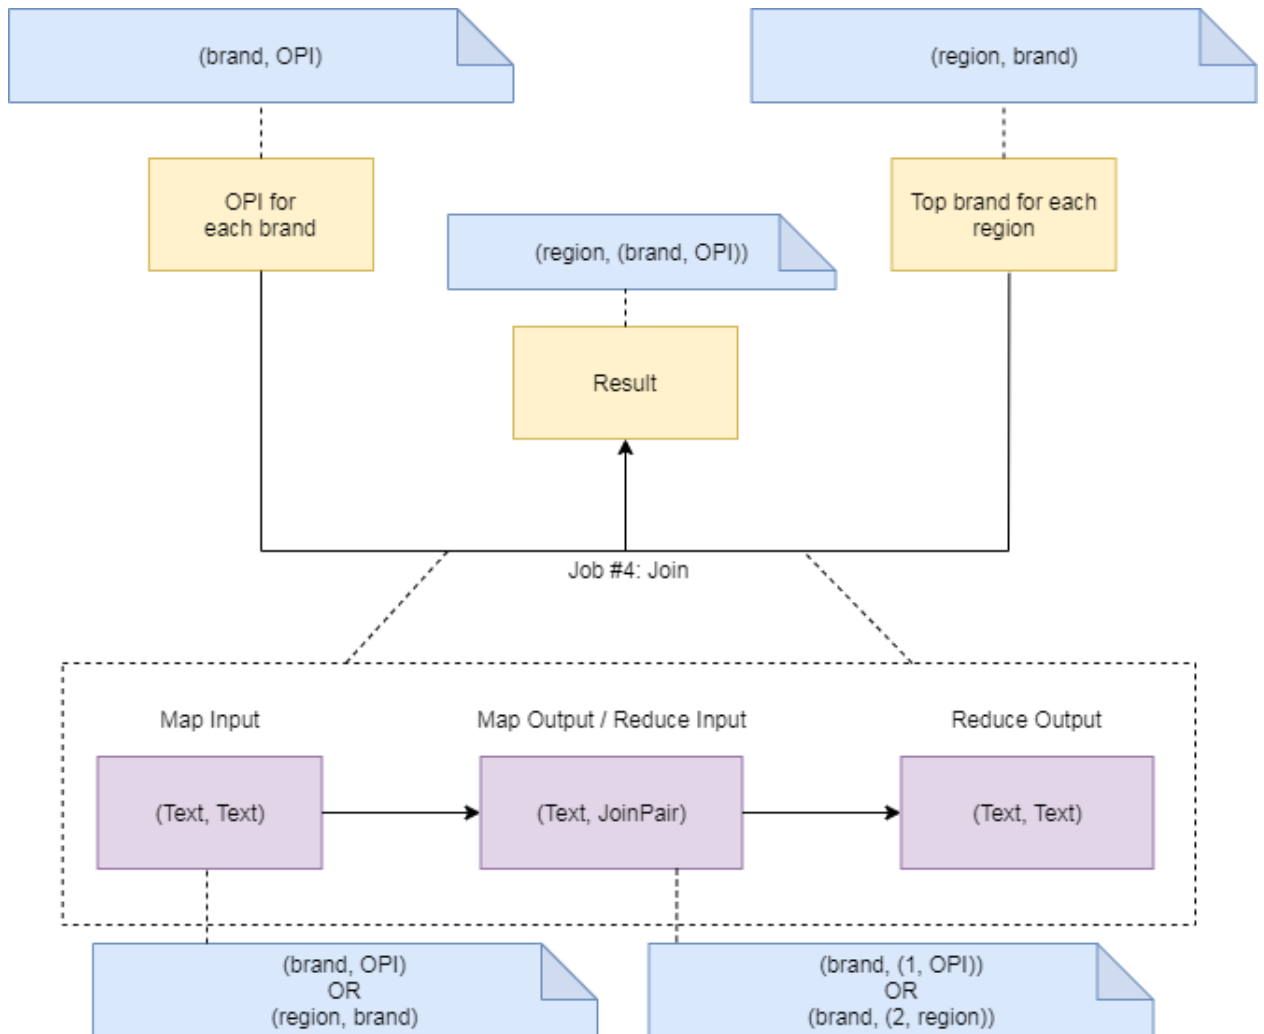


Figure 2.4: Description of “job 3: join” workflow and interfaces.

## 2.2 Performance evaluation

Performance are evaluated by varying the number of reducers, by enabling the use of combiners and by considering different replication factors.

Each casuistry result is the average value upon 3 runs. Results are shown in tables 2.1 and 2.2.

		Reducers				
		1	2	5	10	100
Combiner	No	223	162	168	216	1400
	Yes					

Table 2.1: Performance results expressed as elapsed time in seconds using **replication factor = 1**.

		Reducers				
		1	2	5	10	100
Combiner	No					
	Yes					

Table 2.2: Performance results expressed as elapsed time in seconds using **replication factor = 100**.

# Chapter 3

## Spark

### 3.1 Job description

All jobs specified in section 1.2.2 are executed in the same Spark file. The dataset is loaded as a Spark “DataFrame” and every operation will lean on Spark SQL.

Figure 3.1 shows the complete DAG corresponding to the main job created by Spark. Figure 3.2 shows Spark DAG accompanied with few extra info.


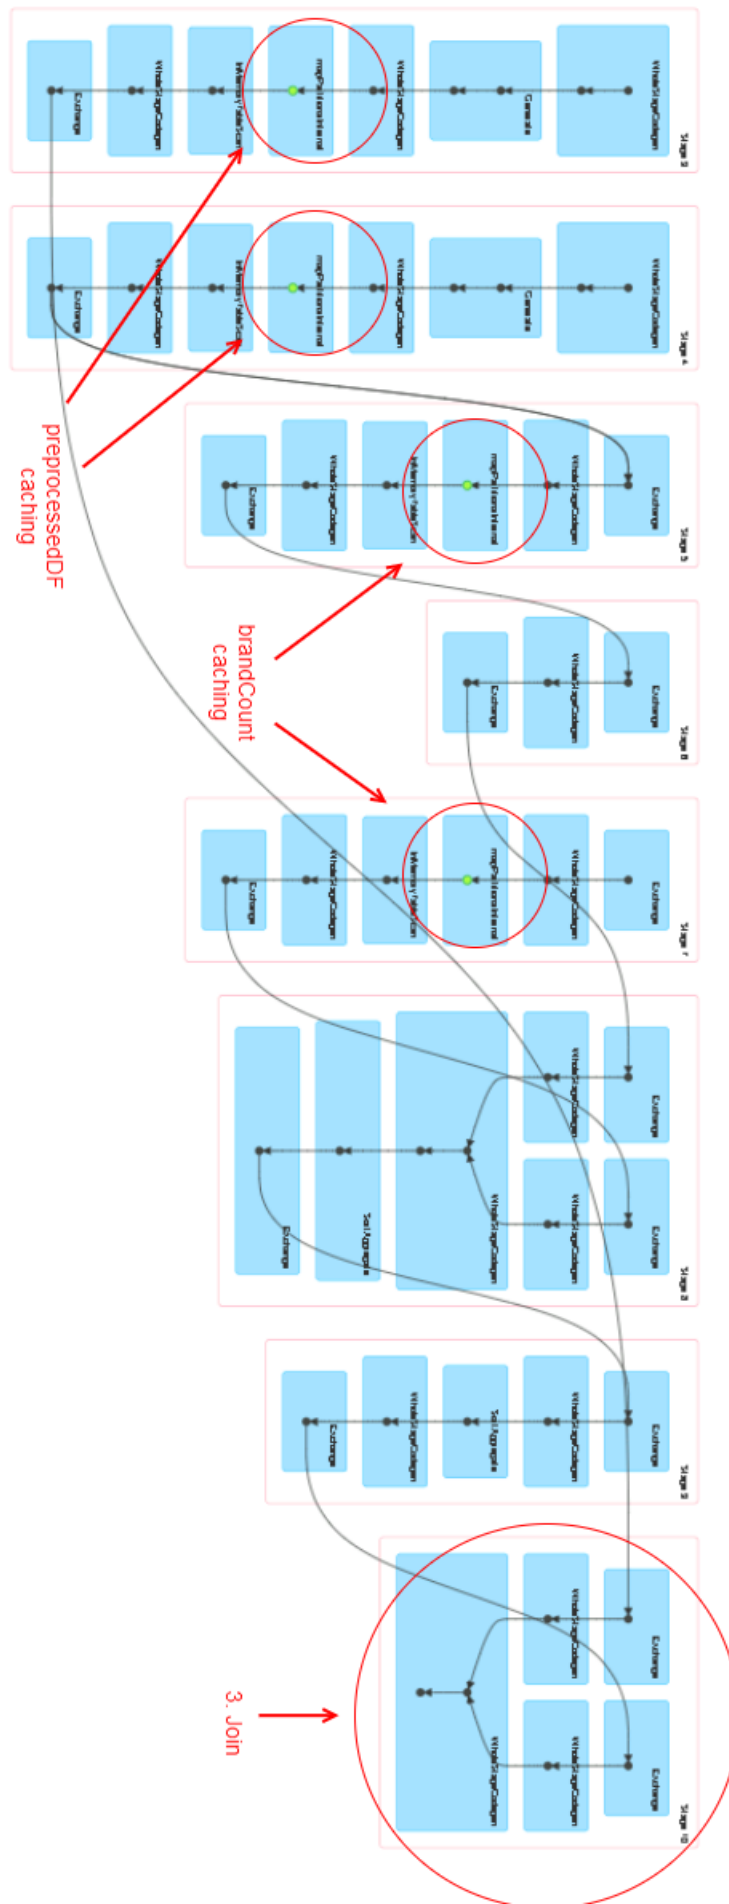


Figure 3.1: Complete Spark DAG.



Since the size of tables used in the two join operations is not excessive, the usage of the broadcast feature may be used on one table. Such an operation makes Spark split the job in three pipelined different jobs (few stages are skipped though).

DAGs are reported in figures 3.3, 3.4 and 3.5.

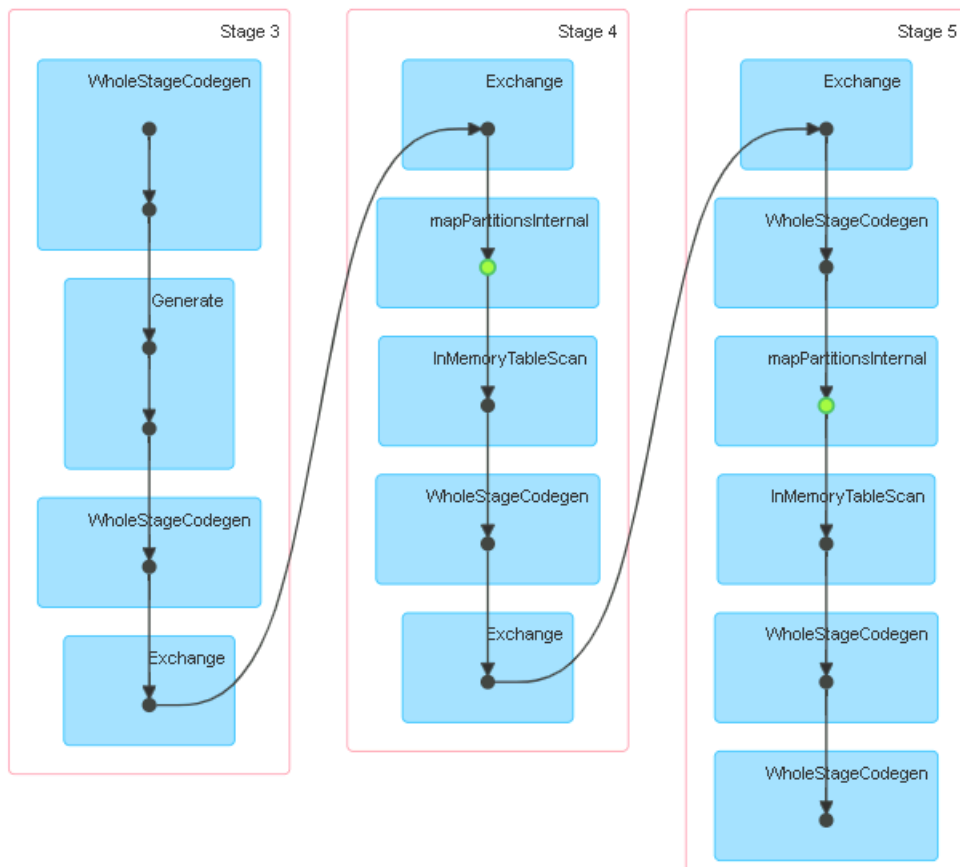


Figure 3.3: DAG for first job created by Spark when using broadcast on join.



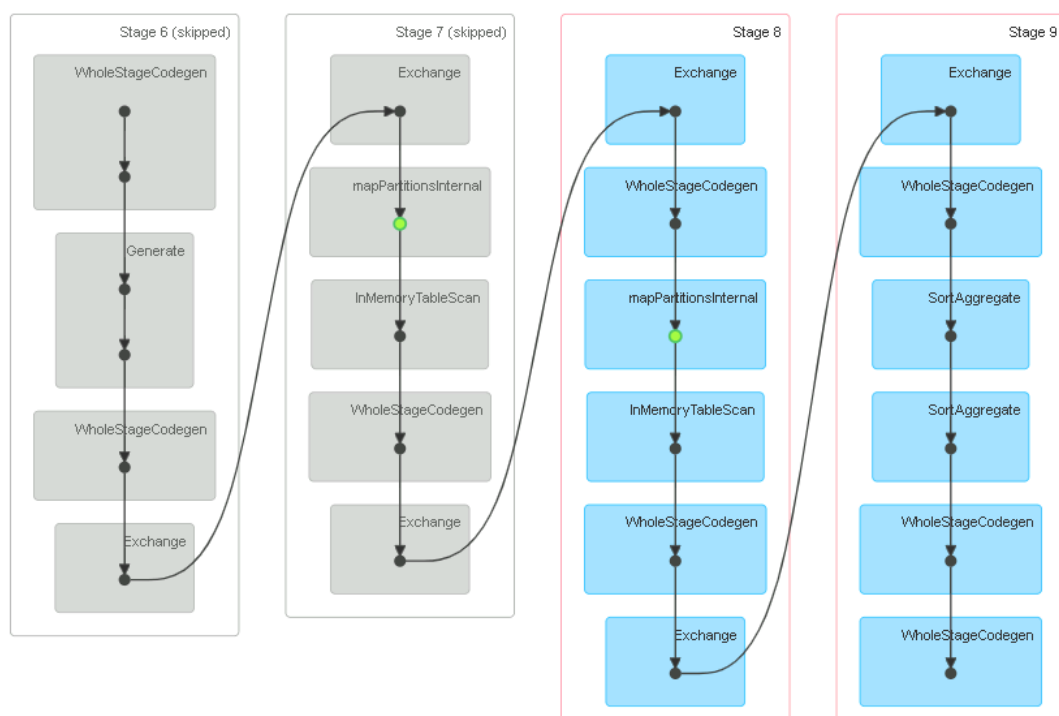


Figure 3.4: DAG for second job created by Spark when using broadcast on join.

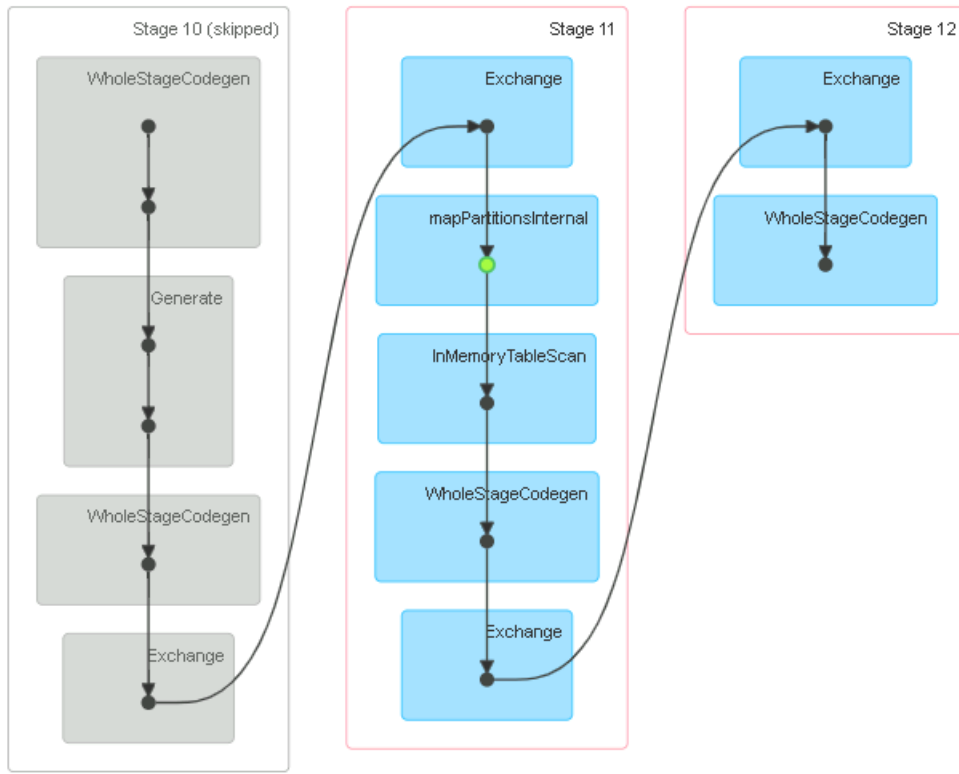


Figure 3.5: DAG for third job created by Spark when using broadcast on join.

## 3.2 Performance evaluation

Performance are evaluated by varying the number of executors and cores per executor. Let's consider cluster characteristics (10 nodes, 4 cores per node) and given constraints (usage of 8 CPUs at most). Each casuistry result is the average value upon 3 runs. Results are shown in table 3.1.

As we can notice, results confirm common best practice rules: best performances among all more promising combinations are obtained using the maximum number of resources ( 8 CPUs); moreover, the whole application speeds up when assigning more than one core to each executor.

		Executors				
		1	2	3	4	8
Cores	1					88
per	2			94	79	
executor	3		96			

Table 3.1: Performance results expressed as elapsed time in seconds using **replication factor = 100** for most promising combinations.



## Chapter 4

## Conclusions

